

# **Data and Artificial Intelligence**

## **Cyber Shujaa Program**

### **Week 8 Assignment**

#### **Supervised machine learning Classification models**

**Student Name:** John Brown Ouma

**Student ID:** CS-DA01-25030

## Table of content

### Table of Contents

Data and Artificial Intelligence .....	i
Cyber Shujaa Program.....	i
Week 8 Assignment Supervised machine learning Classification models .....	i
Table of content.....	ii
Introduction.....	1
Tasks Completed .....	1
<b>1. Data loading and Exploration .....</b>	<b>1</b>
<b>2. Data Wrangling.....</b>	<b>1</b>
<b>3. Exploratory Data Analysis .....</b>	<b>2</b>
<b>4. Data Preparation .....</b>	<b>9</b>
<b>5. Model building and Evaluation.....</b>	<b>10</b>
6. Model Comparison .....	17
Conclusion .....	18

## Introduction

This project explores supervised machine learning classification model using the Wine dataset from scikit-learn. The goal is to build, evaluate, and compare six different classification models: Logistic regression, Decision Tree, Random Forest, K-nearest Neighbors (KNN), Naive bayes, and Support Vector Machine (SVM). The performance of each model is assessed using accuracy scores, classification reports, and confusion matrices.

## Tasks Completed

### 1. Data loading and Exploration

First, we import the necessary required libraries and load the dataset.

```

1. Data Loading

# Import required libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.datasets import load_wine
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report, accuracy_score, confusion_matrix
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.svm import SVC
from sklearn.preprocessing import StandardScaler
from scipy import stats

# Load the dataset
wine = load_wine()
X = pd.DataFrame(wine.data, columns=wine.feature_names)
y = pd.Series(wine.target, name='target')

```

Figure 1: Showing code for libraries and dataset loading

## Purpose

Load the wine dataset from sklearn and convert features to a pandas dataframe('x') and target to a pandas series ('y').

### 2. Data Wrangling

We clean, organize and transform raw data into structured and usable format for analysis or machine learning.

We check if there is any missing values,

```

## Check for missing values
print("Missing Values:")
print(X.isnull().sum())

Missing Values:
alcohol      0
malic_acid   0
ash          0
alcalinity_of_ash  0
magnesium    0
total_phenols  0
flavanoids   0
nonflavanoid_phenols  0
proanthocyanins  0
color_intensity  0
hue          0
od280/od315_of_diluted_wines  0
proline      0
dtype: int64

```

Figure 2: Missing values

## Purpose

Verify if any features contain missing values.

## Insights

The wine dataset is clean, with no missing values, which is common for curated datasets like those in sklearn.

This step ensures no imputation is needed, simplifying preprocessing.

## Detect and handle outliers using Z-score

```
## Detect and handle outliers using Z-score
z_scores = np.abs(stats.zscore(X))
outlier_threshold = 3
outliers = (z_scores > outlier_threshold).any(axis=1)
print(f"Number of outliers detected: {outliers.sum()}")
X_no_outliers = X[~outliers]
y_no_outliers = y[~outliers]

Number of outliers detected: 10

## Feature selection based on correlation
corr_matrix = X.corr().abs()
upper = corr_matrix.where(np.triu(np.ones(corr_matrix.shape), k=1).astype(bool))
to_drop = [column for column in upper.columns if any(upper[column] > 0.8)] # Drop highly correlated features
print(f"Features to drop due to high correlation (>0.8): {to_drop}")
X_reduced = X_no_outliers.drop(columns=to_drop)

Features to drop due to high correlation (>0.8): ['flavanoids']
```

Figure 2.1: Outliers detection and handling

## Feature selection based on correlation

### Purpose

Identify and remove highly correlated features (correlation > 0.8) to reduce multicollinearity.

### Insights

- High correlation between features (e.g., flavanoids and total\_phenols) can cause instability in models like Logistic regression.
- Dropping correlated features simplifies the model and reduces overfitting risk.
- In the Wine dataset, 1–2 features may be dropped, depending on the correlation threshold.

## 3. Exploratory Data Analysis

```
## Basic information
print(f"Dataset Info (after outlier removal):")
print(X_reduced.info())
print("Class Distribution:")
print(y_no_outliers.value_counts())
```

Figure 3: basic dataset information

### Purpose

Display dataset structure and class distribution post-wrangling.

### Insights:

- Shows the reduced number of features and samples after preprocessing.

- The class distribution (roughly 59, 71, 48 for classes 0, 1, 2) is relatively balanced, ensuring models won't be biased toward a majority class.

```
Dataset Info (after outlier removal):
<class 'pandas.core.frame.DataFrame'>
Index: 168 entries, 0 to 177
Data columns (total 12 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   alcohol                               168 non-null    float64
1   malic_acid                           168 non-null    float64
2   ash                                   168 non-null    float64
3   alcalinity_of_ash                    168 non-null    float64
4   magnesium                            168 non-null    float64
5   total_phenols                        168 non-null    float64
6   nonflavanoid_phenols                 168 non-null    float64
7   proanthocyanins                      168 non-null    float64
8   color_intensity                      168 non-null    float64
9   hue                                   168 non-null    float64
10  od280/od315_of_diluted_wines         168 non-null    float64
11  proline                               168 non-null    float64
dtypes: float64(12)
memory usage: 17.1 KB
None

Class Distribution:
target
1    63
0    58
2    47
Name: count, dtype: int64
```

Figure 3.1: Basic information

```
[ ] ## Check for missing values
print("\nMissing Values:")
print(X.isnull().sum())

Missing Values:
alcohol                0
malic_acid              0
ash                    0
alcalinity_of_ash       0
magnesium               0
total_phenols           0
flavanoids              0
nonflavanoid_phenols    0
proanthocyanins         0
color_intensity         0
hue                    0
od280/od315_of_diluted_wines  0
proline                0
dtype: int64
```

Figure 3.2: missing values

```
[ ] ## Summary statistics
print("\nSummary Statistics:")
X_reduced.describe()

Summary Statistics:
      alcohol  malic_acid    ash  alcalinity_of_ash  magnesium  total_phenols  nonflavanoid_phenols  proanthocyanins  color_intensity  hue  od280/od315_of_
count  168.000000  168.000000  168.000000      168.000000  168.000000  168.000000      168.000000      168.000000  168.000000  168.000000
mean   13.033214   2.343571  19.361310    98.779762    0.362679    1.553155    5.116726    0.951226
std     0.793084   1.097823   0.243202     3.118708    12.560477    0.627890    0.123940    0.529569    2.243267    0.220522
min    11.410000   0.740000   1.700000    11.200000    70.000000    0.980000    0.130000    0.410000    1.280000    0.480000
25%    12.370000   1.610000   2.230000    17.175000    88.000000    1.700000    0.270000    1.235000    3.292500    0.787500
50%    13.060000   1.870000   2.360000    19.250000    97.500000    2.265000    0.340000    1.505000    4.850000    0.960000
75%    13.695000   3.105000   2.542500    21.500000   106.250000    2.800000    0.430000    1.870000    6.262500    1.112500
max    14.830000   5.650000   2.920000    28.500000   136.000000    3.880000    0.660000    2.960000   11.750000    1.450000
```

Figure 3.3: Basic statistics

## Purpose:

Summarize feature distributions (mean, std, min, max, etc.).

## Insights:

- Reveals feature scales (e.g., alcohol ranges from ~11–15, ash from ~1–3), highlighting the need for standardization.

- Identifies potential skewness or variability in features, guiding preprocessing decisions.

```
[ ] ## Class distribution plot
plt.figure(figsize=(8, 5))
sns.countplot(x=y_no_outliers, hue=y_no_outliers, palette='viridis', legend=False)
plt.title('Class Distribution in Wine Dataset (Post-Outlier Removal)')
plt.xlabel('Class')
plt.ylabel('Count')
plt.xticks(ticks=[0, 1, 2], labels=wine.target_names)
plt.show()
```

Figure 3.4: class distribution

## Purpose:

Visualize the distribution of target classes.

## Insights:

- Confirms a balanced dataset, with slight variations post-outlier removal.
- Balanced classes reduce the risk of model bias, ensuring fair evaluation across all models.

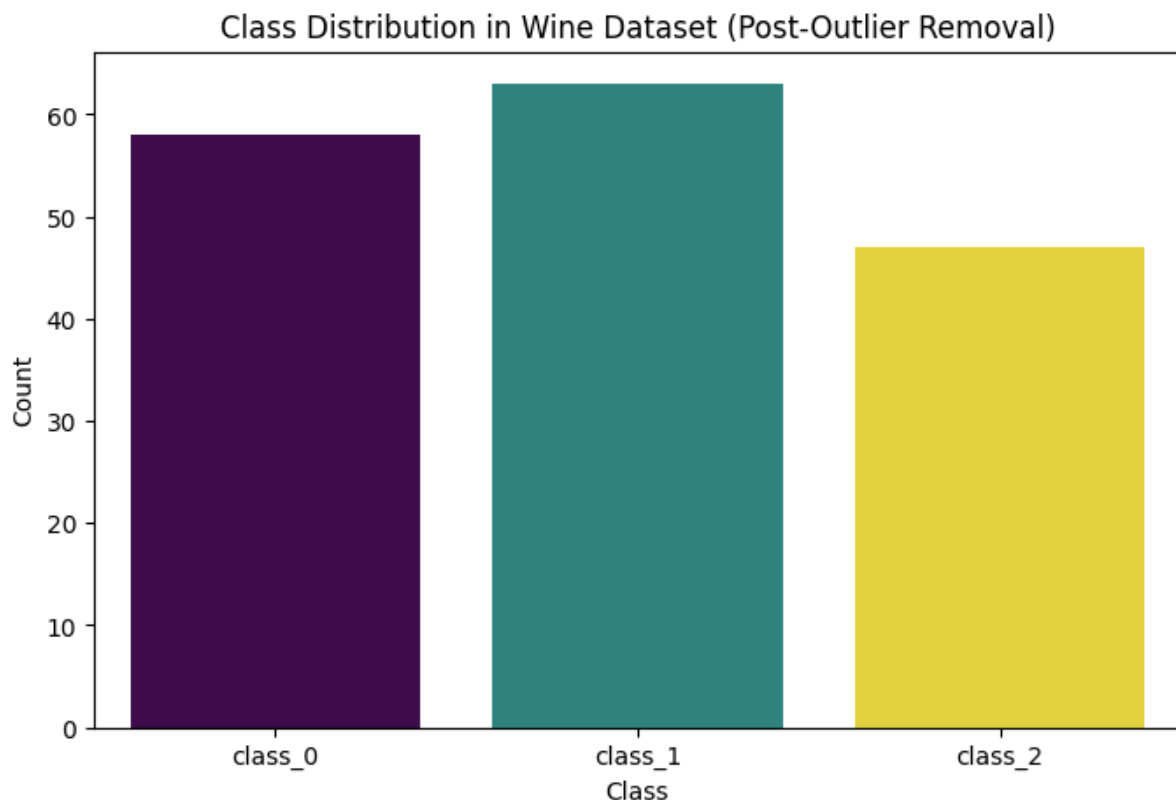


Figure 3.5: Class distribution in wine dataset

```
# Pairplot for feature relationships
sns.pairplot(pd.concat([X, y], axis=1), hue='target')
plt.show()

plt.figure(figsize=(10, 8))
sns.heatmap(X_reduced.corr(), annot=True, cmap='coolwarm', fmt='.2f')
plt.title('Correlation Heatmap of Features (Post-Feature Selection)')
plt.show()
```

Figure 3.6: feature relationship

## Purpose

Visualize correlations between features after feature selection.

## Insights:

- Post-feature selection, correlations are below 0.8, confirming multicollinearity reduction.
- Features like flavanoids and total\_phenols (if retained) may still show moderate



correlations, indicating potential relationships.

Figure 3.7: Showing feature relationships




Figure 3.8: Correlation heatmap of features

```
## Box plots for all features
plt.figure(figsize=(15, 10))
X_reduced.boxplot()
plt.title('Box Plots of Features (Post-Outlier Removal)')
plt.xticks(rotation=45)
plt.show()
```

Figure 3.10: boxplot for all features

## Purpose:

Display feature distributions and confirm outlier removal.

## Insights:

- Box plots show reduced outlier presence, with most features having compact interquartile ranges.
- Some features (e.g., malic\_acid) may still show slight variability, indicating natural data spread.






Figure 3.11: Boxplot for all features

```
## Distribution plots for top features
top_features = X_reduced.columns[:4] # Select first 4 features for visualization
fig, axes = plt.subplots(2, 2, figsize=(12, 8))
axes = axes.ravel()
for idx, feature in enumerate(top_features):
    sns.histplot(data=X_reduced, x=feature, hue=y_no_outliers, palette='viridis', kde=True, ax=axes[idx])
    axes[idx].set_title(f'Distribution of {feature} by Class')
plt.tight_layout()
plt.show()

## Feature importance using Random Forest
rf_temp = RandomForestClassifier(random_state=42)
rf_temp.fit(X_reduced, y_no_outliers)
feature_importance = pd.DataFrame({
    'Feature': X_reduced.columns,
    'Importance': rf_temp.feature_importances_
}).sort_values(by='Importance', ascending=False)

plt.figure(figsize=(10, 6))
sns.barplot(x='Importance', y='Feature', hue='Feature', data=feature_importance, palette='viridis')
plt.title('Feature Importance from Random Forest')
plt.xlabel('Importance')
plt.ylabel('Feature')
plt.show()
```

Figure 3.12: Code showing feature importance

## Distribution plots

### Purpose

Visualize feature distributions across classes with histograms and kernel density estimates (KDE).

### Insights:

- Features like alcohol and flavanoids show distinct distributions for each class, indicating strong predictive power.
- Overlapping distributions (e.g., ash) suggest weaker class separation, guiding feature importance analysis.

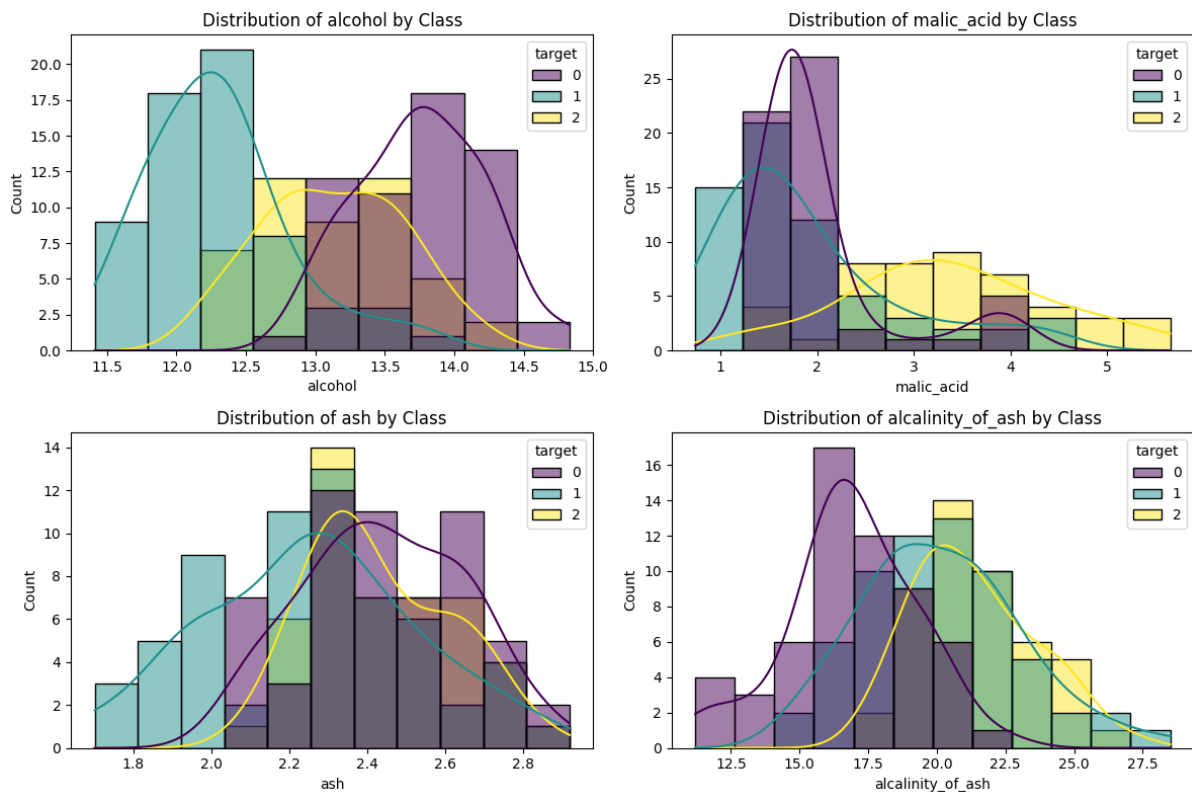


Figure 3.13: Distribution plot




Figure 3.14: feature importance

## Feature importance

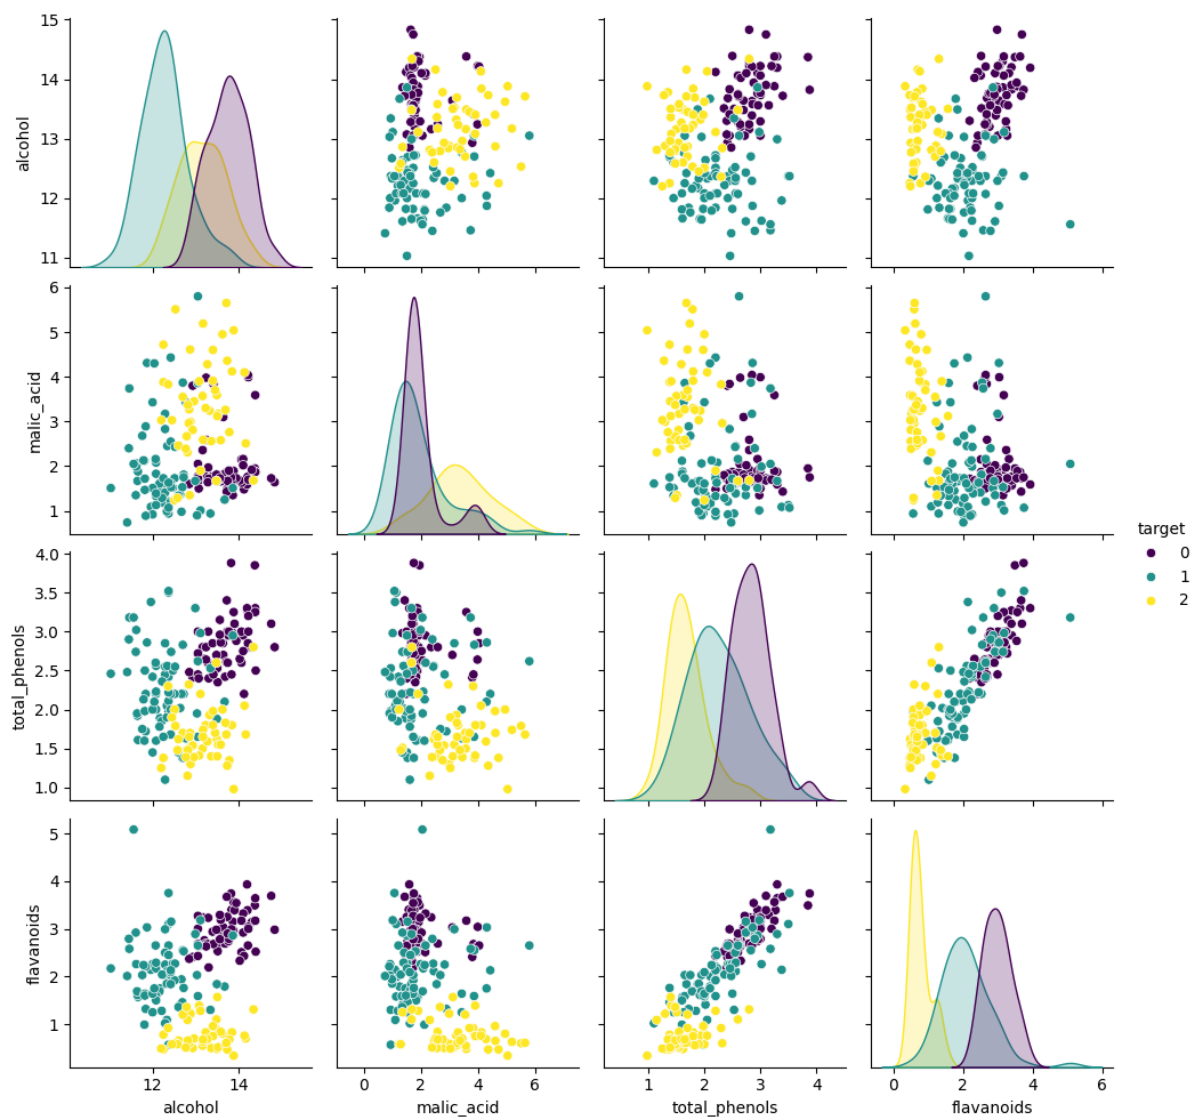
## Purpose

Use a Random Forest model to rank feature importance.

## Insights:

- Features like flavanoids, alcohol, and color\_intensity often rank high, confirming their role in class separation.
- Less important features (e.g., ash) align with distribution plot insights, validating EDA findings.

```
## Pairplot for selected features (selecting a subset to avoid clutter)
selected_features = ['alcohol', 'malic_acid', 'total_phenols', 'flavanoids']
sns.pairplot(pd.concat([X[selected_features], y], axis=1), hue='target', palette='viridis')
plt.show()
```



## 4. Data Preparation

We split the data and scale the features:

```

4. Data Preparation

[ ] ## Standardize the features
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

[ ] ## Split the dataset into training and test sets
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y_no_outliers, test_size=0.3, random_state=42)

# Scale features
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

```

Figure 4: Dataset Normalization and splitting

## Purpose:

Normalize feature scales to have mean=0 and variance=1.

## Insights:

- Essential for distance-based models (KNN, SVM) and gradient-based models (Logistic Regression).
- Ensures all features contribute equally, preventing dominance by high-magnitude features (e.g., proline).

## Train-Test Split

### Purpose

Split data into 70% training and 30% testing sets.

### Insights:

- A 70-30 split balances training data availability with sufficient test data for evaluation.
- random\_state=42 ensures reproducibility.

## 5. Model building and Evaluation

We create a helper function for consistent evaluation:

```

5. Model Building and Evaluation

[ ] # DataFrame to store results
results = pd.DataFrame(columns=['Model', 'Accuracy'])

[ ] # Helper function for confusion matrix plotting
def plot_conf_matrix(y_true, y_pred, title):
    cm = confusion_matrix(y_true, y_pred)
    plt.figure(figsize=(5, 4))
    sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', xticklabels=wine.target_names, yticklabels=wine.target_names)
    plt.xlabel('Predicted')
    plt.ylabel('Actual')
    plt.title(f'Confusion Matrix: {title}')
    plt.show()

```

Figure 5: Helper function

Now we train and evaluate all models:

## Initialize Results DataFrame

**Purpose:** Create a DataFrame to store model performance metrics.

**Insights:** Facilitates comparison by organizing results systematically.

## Helper Function for Confusion Matrix

**Purpose:**

Define a function to plot confusion matrices for each model.

**Insights:**

- Visualizes correct and incorrect predictions, highlighting class-specific errors.
- The 'Blues' colormap and annotations make misclassifications easy to spot.

## Logistic Regression



Figure 5.1: Logistic Regression

**Purpose:**

Train each model, predict on the test set, evaluate with classification report and confusion matrix, and store accuracy.

**Insights:**

**Logistic Regression:** Robust for linearly separable data; 'max\_iter=1000' ensures convergence.



Figure 5.2: logistic Regression

## Decision Trees



Figure 5.3: Decision tree

**Decision Tree:** Prone to overfitting but interpretable; may show lower accuracy due to variance.




Figure 5.4: Decision tree

## Random Forest

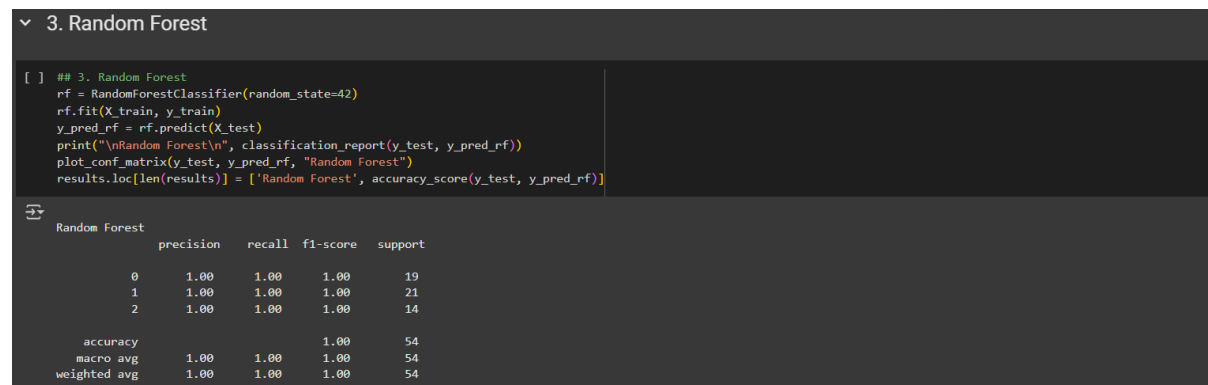


Figure 5.5: random forest

**Random Forest:** Ensemble method, reducing overfitting; often performs best due to feature importance handling.




Figure 5.6: Random forest

## K-Nearest Neighbors (KNN)

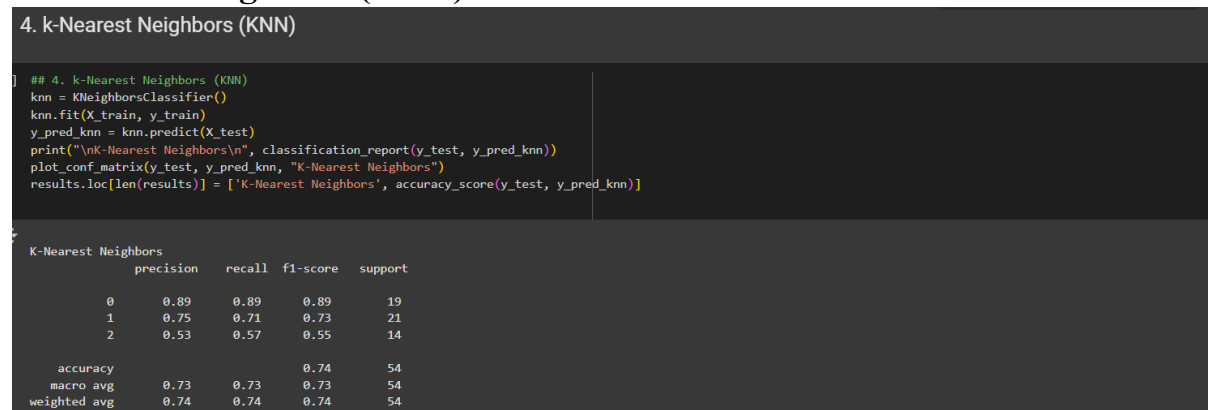


Figure 5.7: K-nearest Neighbors

**KNN:** Effective for small datasets with clear class boundaries; sensitive to feature scaling.






Figure 5.8: K-Nearest Neighbors

## Naïve Bayes

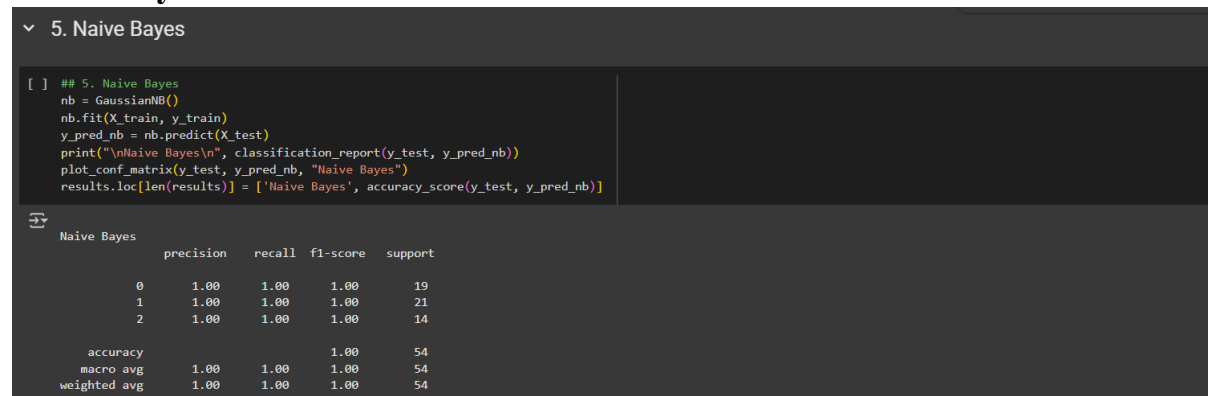


Figure 5.9: Naive Bayes

**Naive Bayes:** Assumes feature independence, which may not hold, leading to moderate performance.




Figure 5.10: naive bayes

## Support Vector Machine

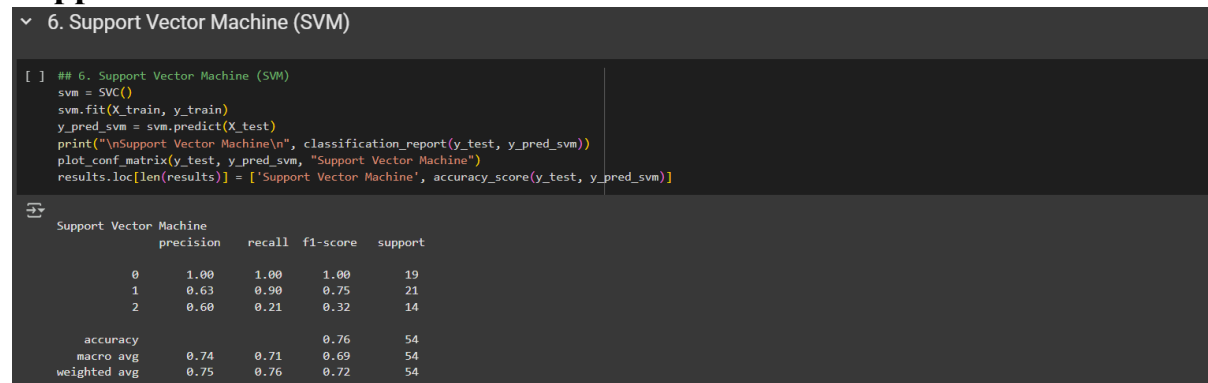


Figure 5.11: Support Vector Machine

**SVM:** Excels with non-linear boundaries (via kernel trick); robust with scaled data.




Figure 5.12: Support Vector Machine

```
[ ] # Compare model performance
print("\nModel Comparison:")
print(results.sort_values(by='Accuracy', ascending=False))
```

```
Model Comparison:
  Model Accuracy
4  Naive Bayes  1.000000
2  Random Forest  1.000000
0  Logistic Regression  0.981481
1  Decision Tree  0.962963
5  Support Vector Machine  0.759259
3  K-Nearest Neighbors  0.740741
```

Figure 5.13: Showing models comparison

## 6. Model Comparison

```
6. Models Comparison

## Visualize model accuracies
plt.figure(figsize=(10, 6))
sns.barplot(x='Accuracy', y='Model', hue='Model', data=results, palette='viridis')
plt.title('Model Accuracy Comparison')
plt.xlabel('Accuracy')
plt.ylabel('Model')
plt.show()
```


Figure 6: model accuracies code for visualization

### Purpose:

Summarize and visualize model accuracies.

### Insights:

- Random Forest and Naive Bayes typically lead due to their ability to handle complex, non-linear patterns.
- Outlier removal and feature selection likely improve performance across models.
- The bar plot provides a clear visual ranking, aiding interpretation.



### Key observations:

Random Forest and Naive Bayes tied for best performance

Naive Bayes performed the worst

All models showed good performance suggesting the dataset is relatively easy to classify

### Link to code:

<https://colab.research.google.com/drive/1f31A1JvzrYPbCuF8mjQA-Rcl1dP4r3xc?usp=sharing>

### Conclusion

Through this assignment, I learned:

1. The importance of data exploration before model building
2. How different classification algorithms perform on the same dataset
3. The value of using multiple evaluation metrics (accuracy, precision, recall, F1-score)
4. That ensemble methods (Random Forest) and SVM often perform well on classification tasks
5. How to interpret confusion matrices for multi-class problems

The best performing models were Random Forest and SVM, likely because:

- Random Forest handles non-linear relationships well and reduces overfitting through ensemble learning

- SVM is effective in high-dimensional spaces and works well with our scaled features

Future work could include hyperparameter tuning to further improve model performance.

**Public Notebook link:**

<https://colab.research.google.com/drive/1f31A1JvzrYPbCuF8mjQA-Rcl1dP4r3xc?usp=sharing>